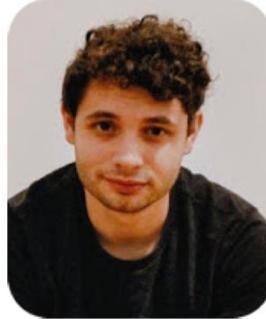


Transformers Learn Shortcuts to Automata

<https://arxiv.org/abs/2210.10749>



Bingbin



Jordan T. Ash



Surbhi Goel



Akshay Krishnamurthy



Cyril Zhang

Algorithmic reasoning

Want to understand the reasoning capabilities & internal mechanisms of self-attention models.

Question: Assume that the variance of the first n natural numbers is 10, and the variance of the first m even natural numbers is 16. Compute $m + n$.

Model output: Let σ^2 be the variance of the first m even natural numbers, and τ^2 be the variance of the first n natural numbers. Then, we have $\tau^2 = \frac{n^2 - 1}{12}$ and $\sigma^2 = \frac{m^2 - 1}{3}$. We solve the system of equations:

$$\frac{n^2 - 1}{12} = 10$$

$$\frac{m^2 - 1}{3} = 16.$$

This translates into the system

$$n^2 = 121$$

$$m^2 = 48 + 1 = 49.$$

Therefore, $n = \boxed{11}$ and $m = \boxed{7}$, so $n + m = \boxed{18}$.

Question: For every $a, b, b \neq a$ prove that

$$\frac{a^2 + b^2}{2} > \left(\frac{a+b}{2}\right)^2.$$

Model output:

$$\begin{aligned} \frac{a^2 + b^2}{2} &> \left(\frac{a+b}{2}\right)^2 \\ \iff \frac{a^2 + b^2}{2} &> \frac{a^2 + b^2 + 2ab}{4} \\ \iff 2a^2 + 2b^2 &> a^2 + b^2 + 2ab \\ \iff a^2 + b^2 &> 2ab \\ \iff a^2 + b^2 - 2ab &> 0 \\ \iff (a - b)^2 &> 0 \end{aligned}$$

which is true, because the square of a real number is positive.



GitHub Copilot

```
1 package main
2
3 type Run struct {
4     Time int // in milliseconds
5     Results string
6     Failed bool
7 }
8
9 // Get average runtime of successful runs in seconds
10 func averageRuntimeInSeconds(runs []Run) float64 {
11     var totalTime int
12     var failedRuns int
13     for _, run := range runs {
14         if run.Failed {
15             failedRuns++
16         } else {
17             totalTime += run.Time
18         }
19     }
20
21     averageRuntime := float64(totalTime) / float64(len(runs) - failedRuns) / 1000
22     return averageRuntime
23 }
```

Code: Codex / Copilot
(Chen et al. 21)

Math: Minerva

(Lewkowycz et al. 22)

Input:
2 9 + 5 7

Target:
<scratch>
2 9 + 5 7 , C: 0
2 + 5 , 6 C: 1 # added 9 + 7 = 6 carry 1
, 8 6 C: 0 # added 2 + 5 + 1 = 8 carry 0
0 8 6
</scratch>
8 6

Scratchpad (Nye et al. 22)

Coin Flip (state tracking)

Q: A coin is heads up. Maybelle flips the coin. Shalonda does not flip the coin. Is the coin still heads up?

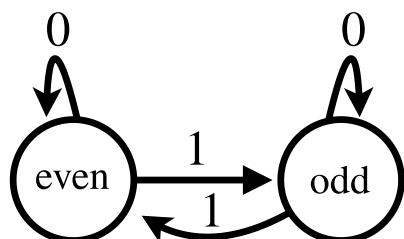
A: The coin was flipped by Maybelle. So the coin was flipped 1 time, which is an odd number. The coin started heads up, so after an odd number of flips, it will be tails up. So the answer is no.

Chain-of-Thought (Wei et al. 21)

Algorithmic reasoning

Coin flip = **parity**

- Input (flip or not): $\Sigma = \{1, 0\}$.
- State (head/even or tail/odd): $Q = \{0, 1\}$.



$$Q = \{\text{even, odd}\}$$
$$\Sigma = \{0, 1\}$$

Coin Flip (state tracking)

Q: A coin is heads up. Maybelle flips the coin. Shalonda does not flip the coin. Is the coin still heads up?

A: The coin was flipped by Maybelle. So the coin was flipped 1 time, which is an odd number. The coin started heads up, so after an odd number of flips, it will be tails up. So the answer is no.

a discrete-time dynamical system

$$\mathcal{A} = (Q, \Sigma, \delta)$$

states inputs transitions

$$q_t = \delta(q_{t-1}, \sigma_t)$$

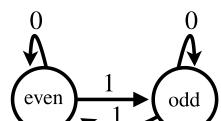
(Semi)Automata

- Semiautomaton \mathcal{A} : a discrete-time dynamical system

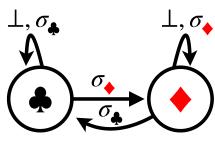
$$\mathcal{A} = (Q, \Sigma, \delta)$$

states inputs transitions

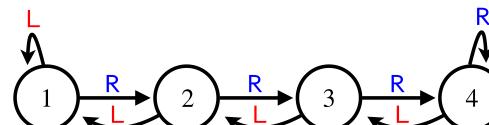
$$q_t = \delta(q_{t-1}, \sigma_t)$$



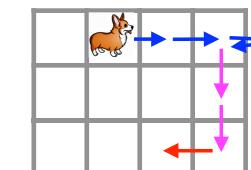
$Q = \{\text{even, odd}\}$
 $\Sigma = \{0, 1\}$
parity counter



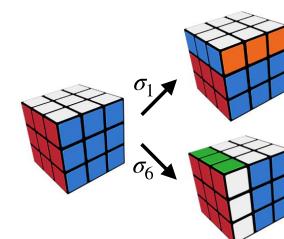
$Q = \{\clubsuit, \diamondsuit\}$
 $\Sigma = \{\sigma_{\clubsuit}, \sigma_{\diamondsuit}, \perp\}$
1-bit memory unit



$Q = \{1, 2, 3, 4\}$
 $\Sigma = \{L, R\}$
1D gridworld



$Q = \{1..3\} \times \{1..4\}$
 $\Sigma = \{\leftarrow, \rightarrow, \uparrow, \downarrow\}$
2D gridworld



$Q = \{54 \text{ stickers}\}$
 $\Sigma = \{6 \text{ face rotations}\}$
Rubik's Cube

- Automaton: $(Q, \Sigma, \delta, \varphi)$, outputs $\tilde{q}_t = \varphi(q_t)$ (acceptance function)

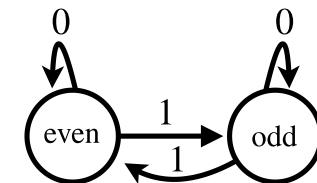
Simulating automata

(Semi)automata: a discrete-time dynamical system

$$\mathcal{A} = (Q, \Sigma, \delta)$$

states inputs transitions

$$q_t = \delta(q_{t-1}, \sigma_t)$$



$$Q = \{\text{even, odd}\}$$

$$\Sigma = \{0, 1\}$$

e.g. **parity counter**

- Task: **simulate \mathcal{A}** : learn a **seq2seq function** for length T .
 - Input: sequence of **tokens** $\sigma_1, \sigma_2, \dots, \sigma_T \in \Sigma$.
 - Output: sequence of **states** $q_1, q_2, \dots, q_T \in Q$.

e.g. parity counter: $\{\sigma_t\} = 01101$, $\{q_t\} = 01001$.

Simulating automata with Transformers



Can Transformers learn automata?

- Automata is *recurrent*; Transformers are *non-recurrent (parallel)* and shallow.

Yes – **Shortcut**: solutions with $o(T)$ sequential “steps/rounds”.

- Def: **longest path** in a computation graph (i.e. at most T).
- e.g. recurrent nets: sequence length.
- e.g. Transformers: number of layers.



Simulating automata with Transformers



Can Transformers learn automata?

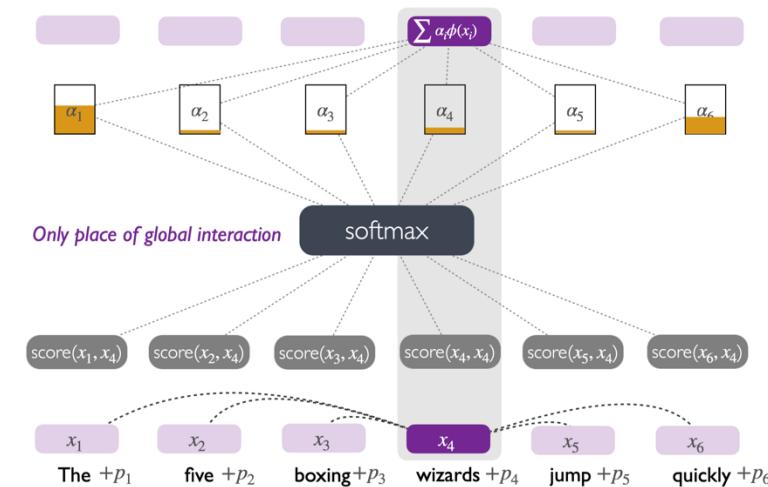
- Automata is *recurrent*; Transformers are *non-recurrent (parallel)* and shallow.

Yes – **Shortcut**: solutions with $o(T)$ sequential “steps/rounds”.

- Def: **longest path** in a computation graph (i.e. at most T).
- e.g. recurrent nets: sequence length.
- e.g. Transformers: number of layers.

Transformer recap

- $\forall i \in [T], x_i^{(l+1)} = \sum_i \alpha_{ij} x_j^{(l)}$;
- $\alpha_{ij} \propto \exp(\langle W_Q x_i^{(l)}, W_K x_j^{(l)} \rangle)$.



Simulating automata with Transformers

🤔 *Can Transformers learn automata?*

- Automata is *recurrent*; Transformers are *non-recurrent (parallel)* and shallow.

Yes – **Shortcut**: solutions with $o(T)$ sequential “steps/rounds”.

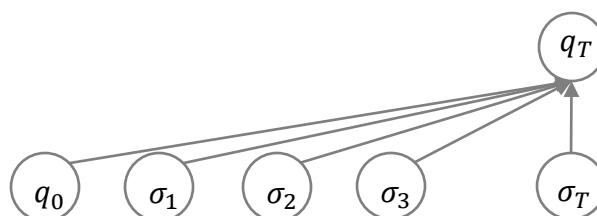
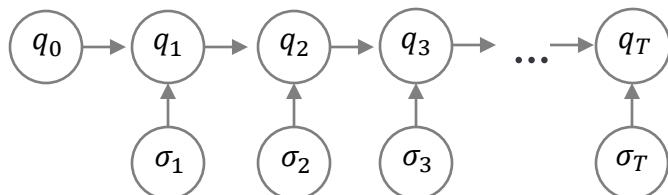
Example: **parity** (prefix sum): $q_t, \forall t \in [n]$,

- Iterative solution:

$$q_t = q_{t-1} \oplus \sigma_t \in \{0,1\}.$$

- Shortcut (parallel) solution:

$$q_t = (\sum_{\tau \leq t} \sigma_\tau) \text{ mod } 2 \in \{0,1\}.$$



Why shortcuts:

- computational advantage;
- Unique & natural to Transformer.

Transformers Learn Shortcuts to Automata

TL;DR: **Shallow** Transformers can simulate $\mathcal{A} = (Q, \Sigma, \delta)$.

- **Theory:** for any length T , Transformers with $o(T)$ layers suffice.
 - $O(\log T)$ -layer simulation for ***all*** \mathcal{A} (also the lower bound for the general case)
 - $O(|Q|^2 \log |Q|)$ -layer simulation for all ***solvable*** \mathcal{A}
 - $O(1)$ -layer simulation for **gridworld** 
- **Empirical study**
 - **Positive:** training shallow models SGD finds shortcuts in practice.
 - **Negative:** they're brittle OOD (distr on input symbols; other lengths).
 - **Fix:** *Scratchpad* training: force a Transformer to learn the recurrent solution.
- **Discussions** 

Theory (as brain teasers)

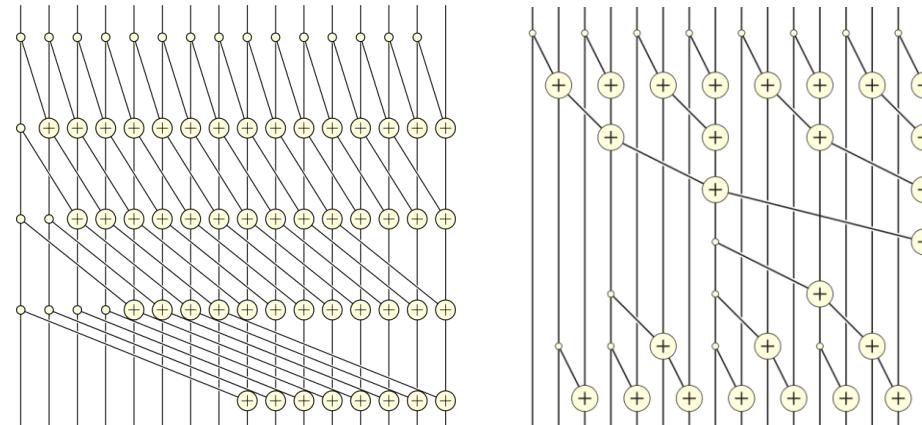
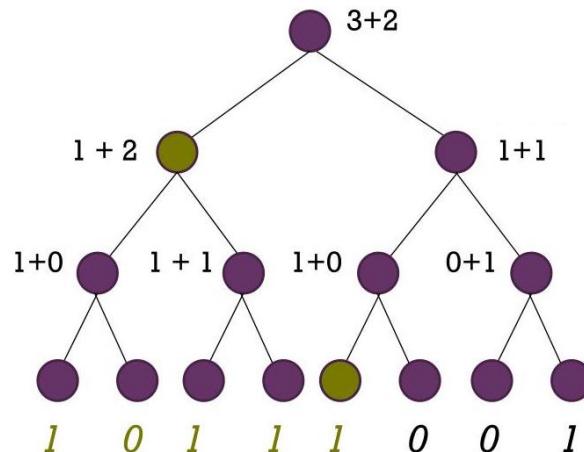
Solutions with $o(T)$ computation depth

Puzzle 1: Do shortcuts exist?

$\mathcal{A} = (Q, \Sigma, \delta)$: (states, inputs, transitions)

Task: for each $t \in [T]$, compute $q_t = (\delta(\cdot, \sigma_t) \circ \dots \circ \delta(\cdot, \sigma_1))(q_0)$.

- Input token $\sigma \rightarrow$ a function $\delta(\cdot, \sigma): Q \rightarrow Q$.
- Function composition (**associative**).

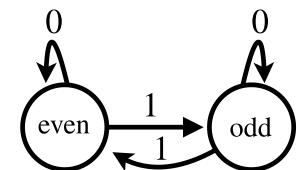


Puzzle 1: Do shortcuts exist?

$\mathcal{A} = (Q, \Sigma, \delta)$: (states, inputs, transitions)

Task: for each $t \in [T]$, compute $q_t = (\delta(\cdot, \sigma_t) \circ \dots \circ \delta(\cdot, \sigma_1))(q_0)$.

- Input token $\sigma \rightarrow$ a function $\delta(\cdot, \sigma): Q \rightarrow Q$.
- Function composition (**associative**) \longleftrightarrow matrix multiplication



$Q = \{\text{even, odd}\}$
 $\Sigma = \{0, 1\}$
parity counter

$$\delta(\cdot, 0) = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \quad (\text{identity})$$

$$\delta(\cdot, 1) = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix} \quad (\text{swap})$$

function composition

$$q_t = (\delta(\cdot, \sigma_t) \circ \dots \circ \delta(\cdot, \sigma_1)) q_0$$

matrix multiplication

$$e_{q_t} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix} \dots \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix} e_{q_0}$$

Representation theory ([Chugtai 23](#)): can explain $O(\log T)$, but not what's coming up.

Can we use $o(\log T)$ layers?

$\mathcal{A} = (Q, \Sigma, \delta)$: (states, inputs, transitions)

We already have positive result on some tasks.

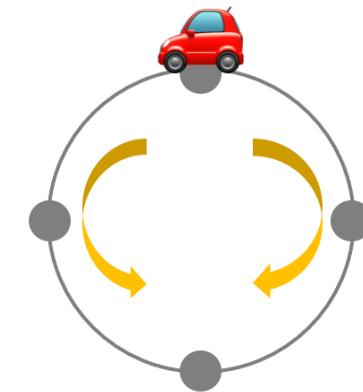
- Parity: $q_t = (\sum_{i \in [t]} \sigma_i) \bmod 2$: only need to **count** #1s.
- Counting suffices, if the function composition is **commutative**.
 - Corollary: any commutative compositions (abelian groups--explained later) can be represented with $O(1)$ layers.

Question: Is there a solution with $o(\log T)$ layers when ***non-commutative***?

Puzzle 2: $\tilde{O}(|Q|^2)$ layers

Reversible car on a circular road:

- $\Sigma = \{D, U\}$ (drive, U-turn), $Q = \{\text{🚗}, \text{🚙}\} \times \{0, 1, 2, 3\}$.
- Consider input $DDDUDDUUD\dots$, starting from $(\text{🚗}, 0)$.
- How to decide the current direction and position?
 - Direction = a parity task on U . (parity: $\{1, -1\} \leftrightarrow \{0, 1\}$)
 - Position = sum of signed counts (sign = parity of U) **mod 4**.



$D \ D \ D \ U \ D \ D \ U \ U \ D$

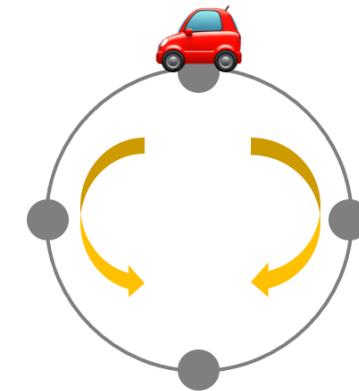
Parity: $1 \ 1 \ 1 \ -1 \ -1 \ -1 \ 1 \ -1 \ -1 \rightarrow \text{🚗}$

Signed counts: $1 \ 1 \ 1 \ 0 \ -1 \ -1 \ 0 \ 0 \ -1 \rightarrow 0$

Solution 2: $\tilde{O}(|Q|^2)$ layers

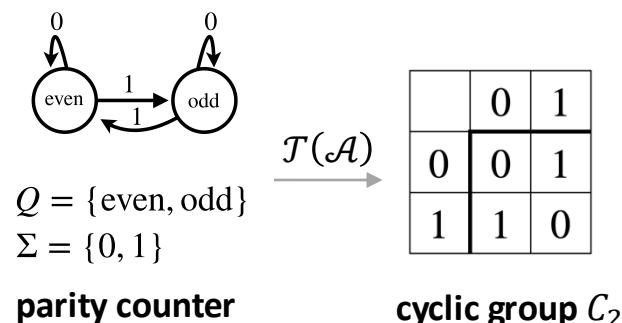
Reversible car on a circular road:

- $\Sigma = \{D, U\}$ (drive, U-turn), $Q = \{0, 1, 2, 3\} \times \{\text{red car}, \text{blue car}\}$.
- Decompose: 1) direction (parity), 2) position (signed sum mod 4).



Is such decomposition always possible? Yes!

Transformation group: $\mathcal{T}(\mathcal{A}) := \{\delta(\cdot, \sigma) : \sigma \in \Sigma\}$ under composition.



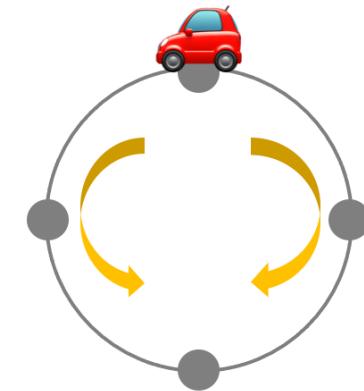
Group G : a set G with operation $G \times G \rightarrow G$.

- Associativity: $(a \cdot b) \cdot c = a \cdot (b \cdot c)$
- Identity: $a \cdot e = e \cdot a = a$
- Inverse: $\forall a \in G, \exists b \in G \text{ s.t. } a \cdot b = b \cdot a = e$

Solution 2: $\tilde{O}(|Q|^2)$ layers

Reversible car on a circular road:

- $\Sigma = \{D, U\}$ (drive, U-turn), $Q = \{0, 1, 2, 3\} \times \{\text{red car}, \text{blue car}\}$.
- Decompose: 1) direction (parity), 2) position (signed sum mod 4).



Is such decomposition always possible? Yes!

Transformation group: $\mathcal{T}(\mathcal{A}) := \{\delta(\cdot, \sigma) : \sigma \in \Sigma\}$ under composition.

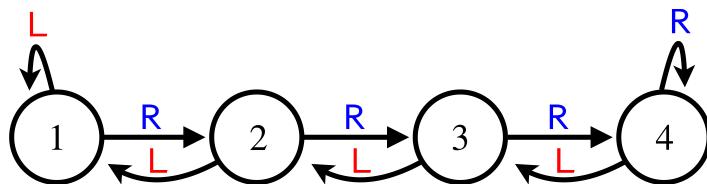
“Prime factorization” for groups: $G \triangleright H_n \triangleright \cdots \triangleright H_1$ (**Jorden & Hölder** [~1880])

- $G \triangleright H$: H is a normal subgroup of G (~factors).
- H_{i+1}/H_i are simple groups (~prime numbers).

What about semigroups?

$\mathcal{T}(\mathcal{A}) := \{\delta(\cdot, \sigma) : \sigma \in \Sigma\}$ under composition

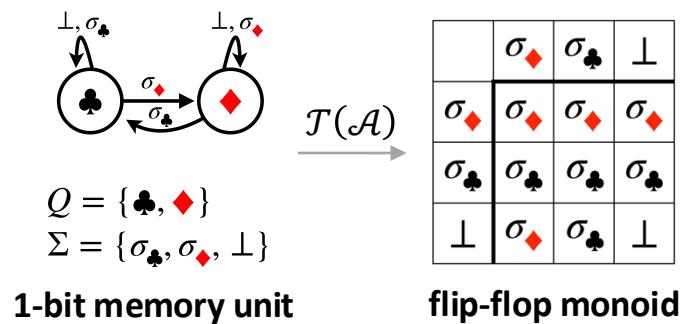
1d gridworld: $\sigma_t = \text{L/R}$ steps, q_t = location in a bounded room.



$$q_0 = 1, \quad q_t = \min(n, \max(1, q_{t-1} + 1))$$

e.g. $\text{LRLLRRRRRL} \mapsto 1, 2, 1, 1, 2, 3, 4, 4, 3, 2$

- $\delta(\cdot, L) = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}$ is not invertible $\rightarrow \mathcal{T}(\mathcal{A})$ is a *semigroup*.



Semigroup G : a generalization of group.

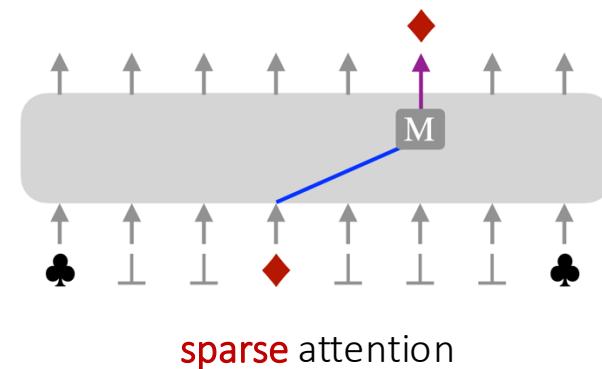
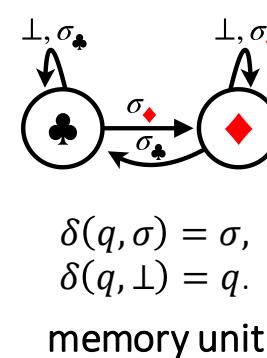
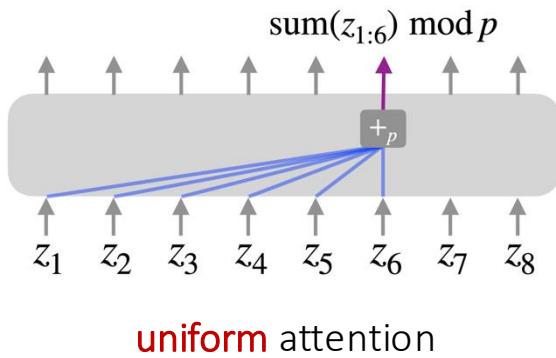
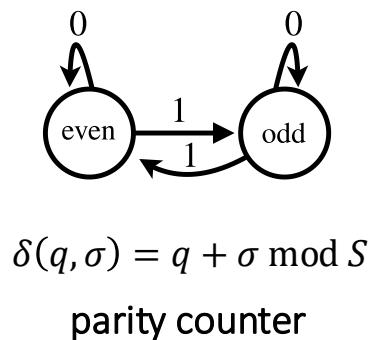
- Associativity.
- (+ Identity: a *monoid*.)

Solution 2: $\tilde{O}(|Q|^2)$ layers

$\mathcal{A} = (Q, \Sigma, \delta)$: (states, inputs, transitions)

1d gridworld: $\sigma_t = \text{L/R}$ steps, q_t = location in a bounded room.

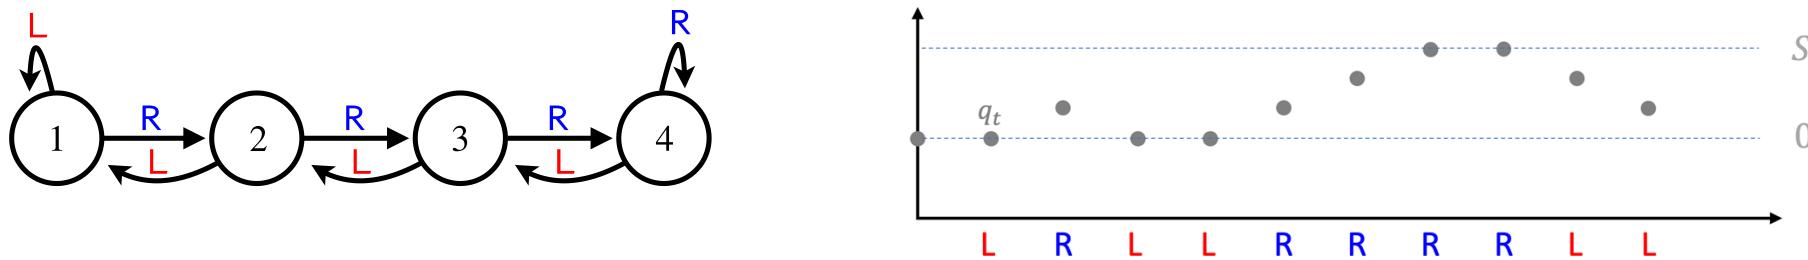
- **Krohn-Rhodes** (Thm 2): solvable semiautomata decomposed into **mod** + **reset**.
 - Solvable: H_{i+1}/H_i is commutative/abelian. [Recall Jorden & Hölder: $G \triangleright H_n \triangleright \dots \triangleright H_1$]



*We do not claim that Transformers learn these decompositions in practice.

Puzzle 3: $O(1)$ layers?

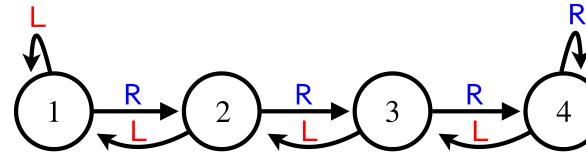
1d gridworld: $\sigma_t = \text{L/R}$ steps, q_t = location in a bounded room.



Puzzle: design a parallel algorithm to compute $\sigma_{1:T} \mapsto q_{1:T}$.

- Hint: *boundary detection*: discard history; easy afterwards (prefix sum).

Solution 3: $O(1)$ layer for

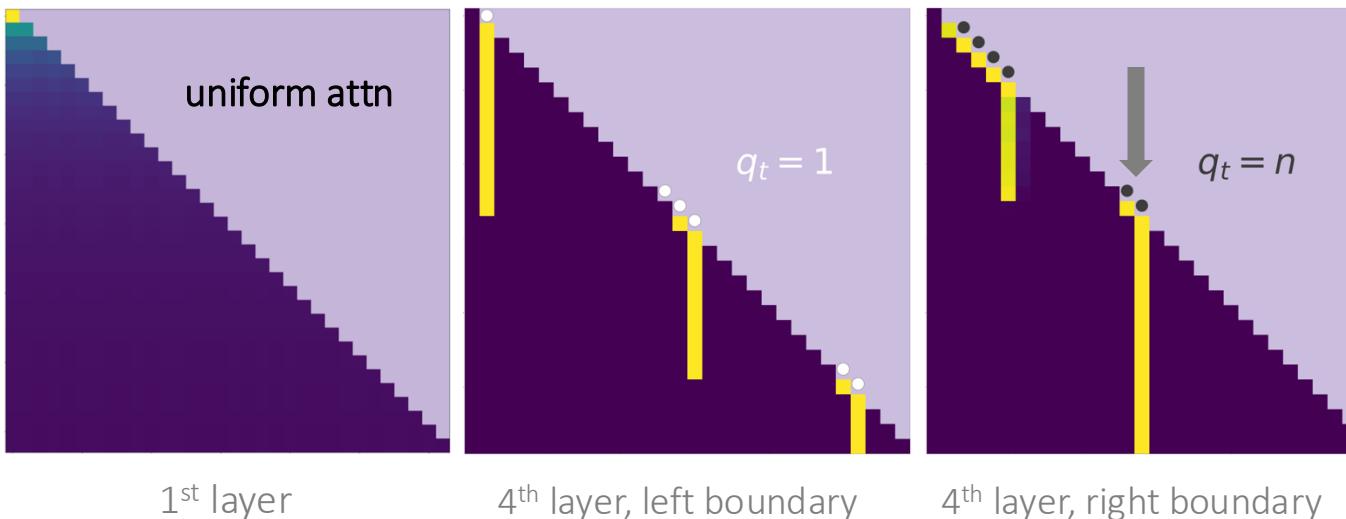


~~"You can only figure out whether you're at a wall if you know q_{t-1} ."~~

- Parallel boundary detection!



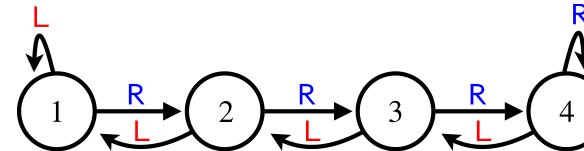
- discard history; easy afterwards (prefix sum).



attention heatmaps
(GPT solved this before we did o o)

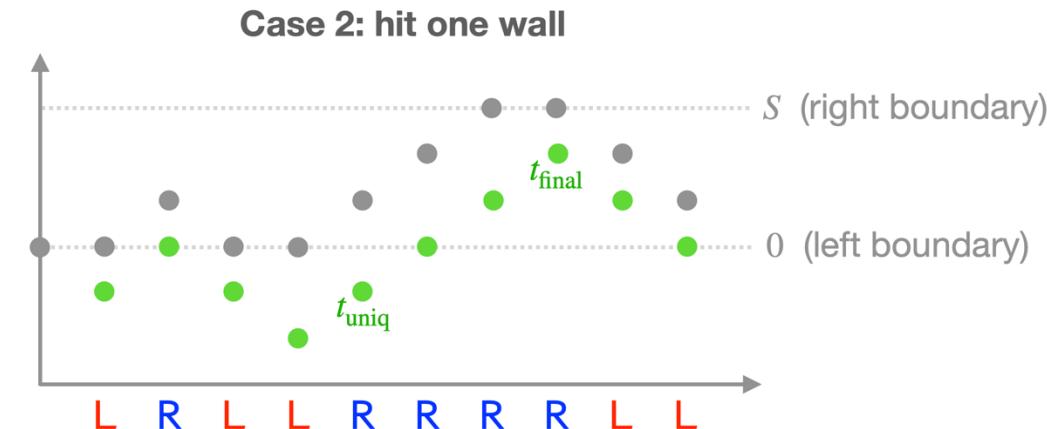
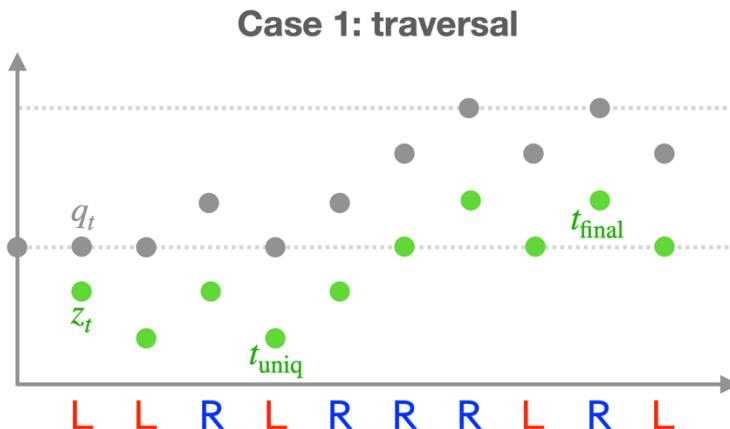
"mechanistic interpretability"
[Nanda & Lieberum '22]

Solution 3: $O(1)$ layer for



- **Parallel boundary detector:**

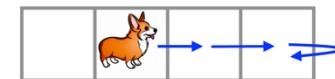
- Compute prefix sums $z_t := \sum \sigma_{1:t}$ (ignoring boundaries);
- At each t , find most recent $t_{\text{uniq}} < t$ such that $z_{t_{\text{uniq}}:t}$ has $n(\# \text{states})$ unique values;
- Then $t_{\text{final}} := \max_{\substack{t_{\text{uniq}} \leq \tau \leq t \\ t_{\text{uniq}} \leq \tau \leq t}} (\text{argmax } z_\tau, \text{ argmin } z_\tau)$ is last boundary collision.



Takeaways

For any length T , Transformers can simulate \mathcal{A} with $o(T)$ depth.

- Theorem 1: $O(\log T)$ -layer simulation for *all* \mathcal{A}
 - Parallel prefix computation (divide and conquer); lower bound for general \mathcal{A} .
- Theorem 2: $O(|Q|^2 \log |Q|)$ -layer simulation for all *solvable* \mathcal{A}
 - Factorization (Krohn-Rhodes).
- Theorem 3: $O(1)$ -layer simulation for *gridworld*
 - Specific \mathcal{A} : even shallower results (e.g. via parallel boundary detector).



Empirical findings

Positive results, challenges

Overview

$$\sigma_{1:T} \rightarrow \text{Transformer} \rightarrow q_{1:T}$$

(Theory)

- Q0: are shallow non-recurrent nets sufficiently expressive? [Yes!]

(Experiments)

- Q1: Can SGD find shortcuts in practice?
- Q2: Shortcomings of shortcuts?
 - Q2.1: Does SGD work with limited supervision?
 - Q2.2: Are shortcuts robust to unseen inputs?
- Q3: Solutions?

(Open question)

- Q4: How to interpret the shortcuts?
- Q5: How to design an architecture to solve reasoning tasks?

SGD works, under ideal supervision

	Transformer depth															
automaton	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Dyck	99.3	100	100	100	100	100	100	100	100	100	100	100	100	100	100	100
Grid ₄	99.9	100	100	100	100	100	100	100	100	100	100	100	100	100	100	100
Grid ₉	92.2	100	100	100	100	100	100	100	100	100	100	100	100	100	100	100
C_2	77.6	99.8	99.9	100	100	99.5	100	99.7	100	100	100	100	100	100	100	100
C_3	54.6	94.6	96.7	99.4	100	100	99.8	100	99.9	100	100	100	100	100	99.8	100
C_4	95.1	92.3	84.2	99.9	99.7	99.9	100	100	100	100	100	100	100	100	100	100
C_5	89.0	99.1	99.9	100	100	100	100	100	100	100	100	100	100	100	100	100
C_6	59.8	98.7	75.5	99.9	99.8	99.9	99.9	100	100	100	99.8	99.9	100	99.8	99.9	99.9
C_7	90.9	95.0	99.9	99.9	100	99.9	100	100	100	100	99.8	100	100	100	100	100
C_8	79.6	96.2	99.8	99.8	99.9	100	99.9	99.9	100	99.4	99.9	99.9	99.9	100	99.9	99.9
C_2^2	90.5	98.8	99.9	100	100	99.9	100	100	99.9	99.9	100	100	100	100	100	100
C_2^3	65.0	77.9	99.9	97.9	100	99.8	98.2	99.9	100	100	91.9	95.9	91.7	90.6	87.5	80.6
D_6	25.4	27.2	47.4	75.2	100	100	100	100	100	100	100	100	100	100	100	100
D_8	45.6	98.0	100	100	100	100	100	100	100	100	100	100	100	100	100	100
Q_8	31.6	49.2	59.6	60.4	73.5	99.3	100	100	100	100	100	100	100	100	100	100
A_4	25.0	35.4	49.1	59.3	62.6	82.3	90.9	98.0	98.0	99.1	99.8	100	99.7	100	100	100
A_5	12.5	23.1	32.5	46.7	71.2	98.8	100	100	100	100	100	100	100	100	100	100
S_4	11.3	17.6	22.0	27.1	37.7	44.8	50.8	72.5	91.3	97.1	97.9	98.7	99.9	100	99.8	99.9
S_5	7.9	11.8	14.6	19.7	26.0	28.4	32.8	51.8	86.3	94.8	90.2	97.2	99.3	99.1	99.9	99.9

Q1: Does SGD on Transformers find shortcuts? *Yes!*

Group-free semiautomata (reset): Dyck, gridworld

- Easy to learn; generalizes Dyck results [Yao et al. '22]

Cyclic groups (mod- n counters): unstable training & o.o.d. eval

- possibly accounts for previous negative results [Bhattamishra et al. '20]
- **open challenge:** improve architectures and training.

Other abelian groups: worse instabilities; higher depth helps training

Harder groups: more results including D_8 , A_5 (non-solvable), S_5 (NC^1 -complete).

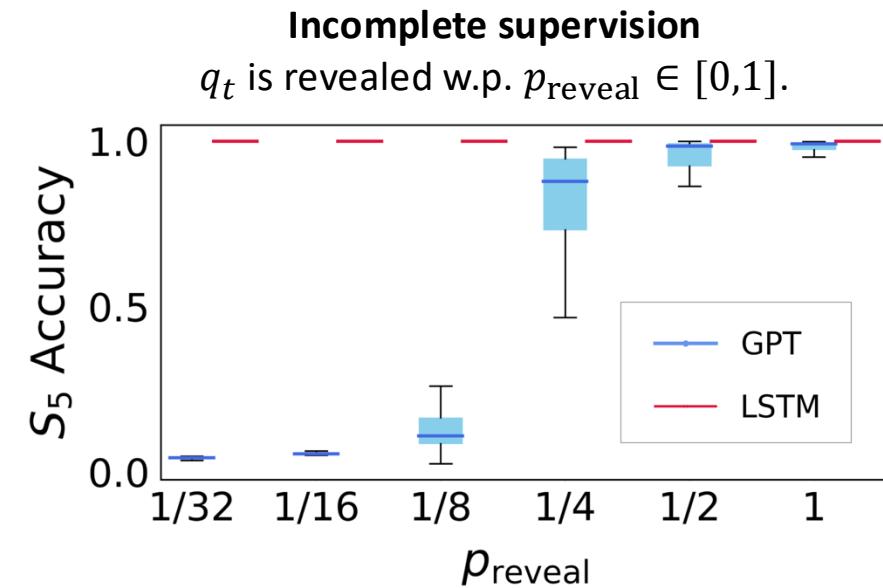
- Groups with deeper factorization requires more layers
- **open challenge 1:** dissect/interpret the learned solutions
- **open challenge 2:** precisely characterize instance complexity

Training robustness: SGD beyond ideal supervision

Q2.1: Does SGD work under limited supervision? *Up to a point.*

- *Setup:* reduce the amount of state info during training; in-distribution eval.

Indirect supervision train & test on a function of q_t .				
Dyck _{4,8}	Grid ₉	S_5	C_4	D_8
stack top	$\mathbb{1}_{\text{boundary}}$	$\pi_{1:t}(1)$	$\mathbb{1}_{0 \bmod 4}$	location
100.0	99.8	99.8	99.7	99.8

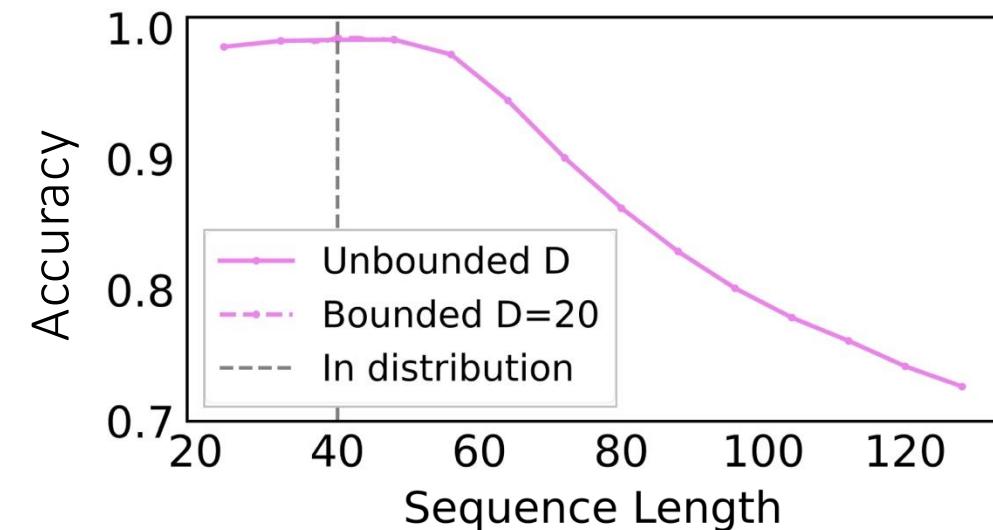
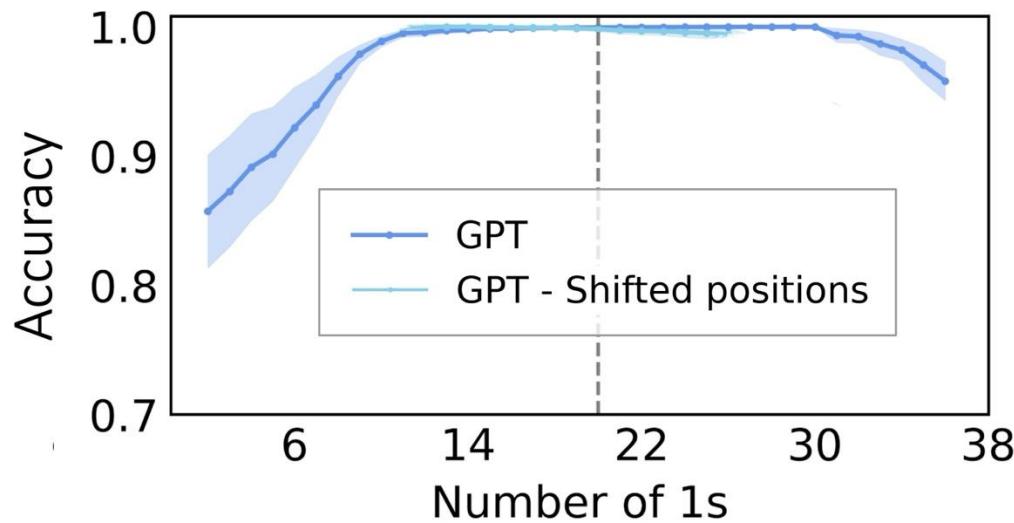


Takeaway: learning is still possible, but not as robust as RNNs (LSTM always 100%).

Testing robustness: hallucinated variables

Q2.2: Are shortcuts robust to unseen inputs? No.

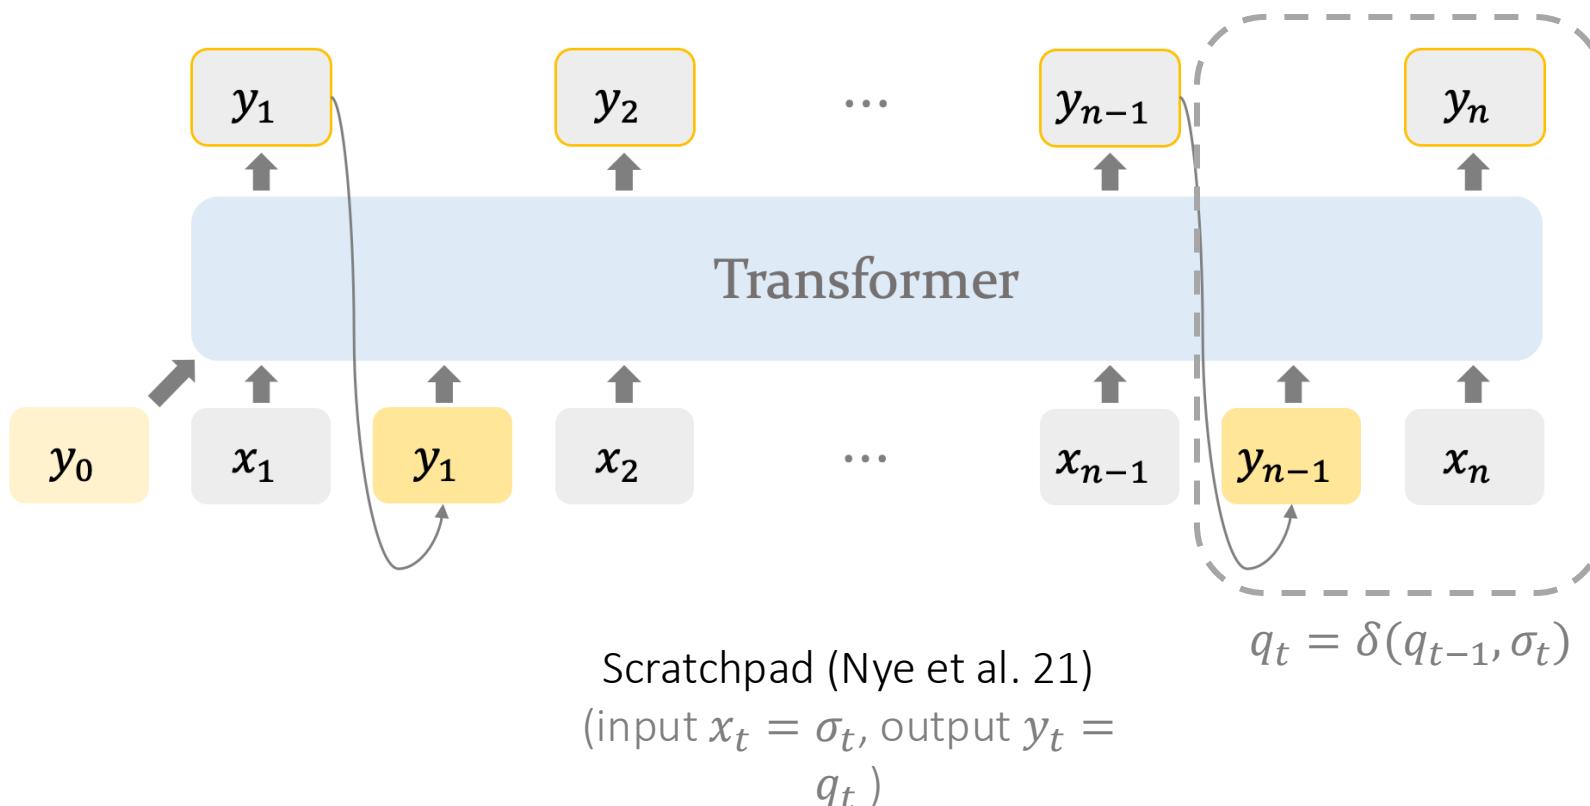
- **Parity (C_2):** ideal training with $p(1) = 0.5$, evaluate under distribution shifts.
 - Shortcut $q_t = (\sum_{i \in [t]} \sigma_i) \bmod 2 \rightarrow (\sum_{i \in [t]} \sigma_i) \approx t/2$ with deviation $\sim \sqrt{t}$.
 - $\bmod 2$ is **memorized** by MLP: fail to generalize to diff $(\sum_{i \in [t]} \sigma_i)$.
- Transformer solutions? Same failure mode \rightarrow maybe same solution.



The recurrent mode of Transformers

Q3: Solutions? Guiding Transformers to learn recurrent solutions.

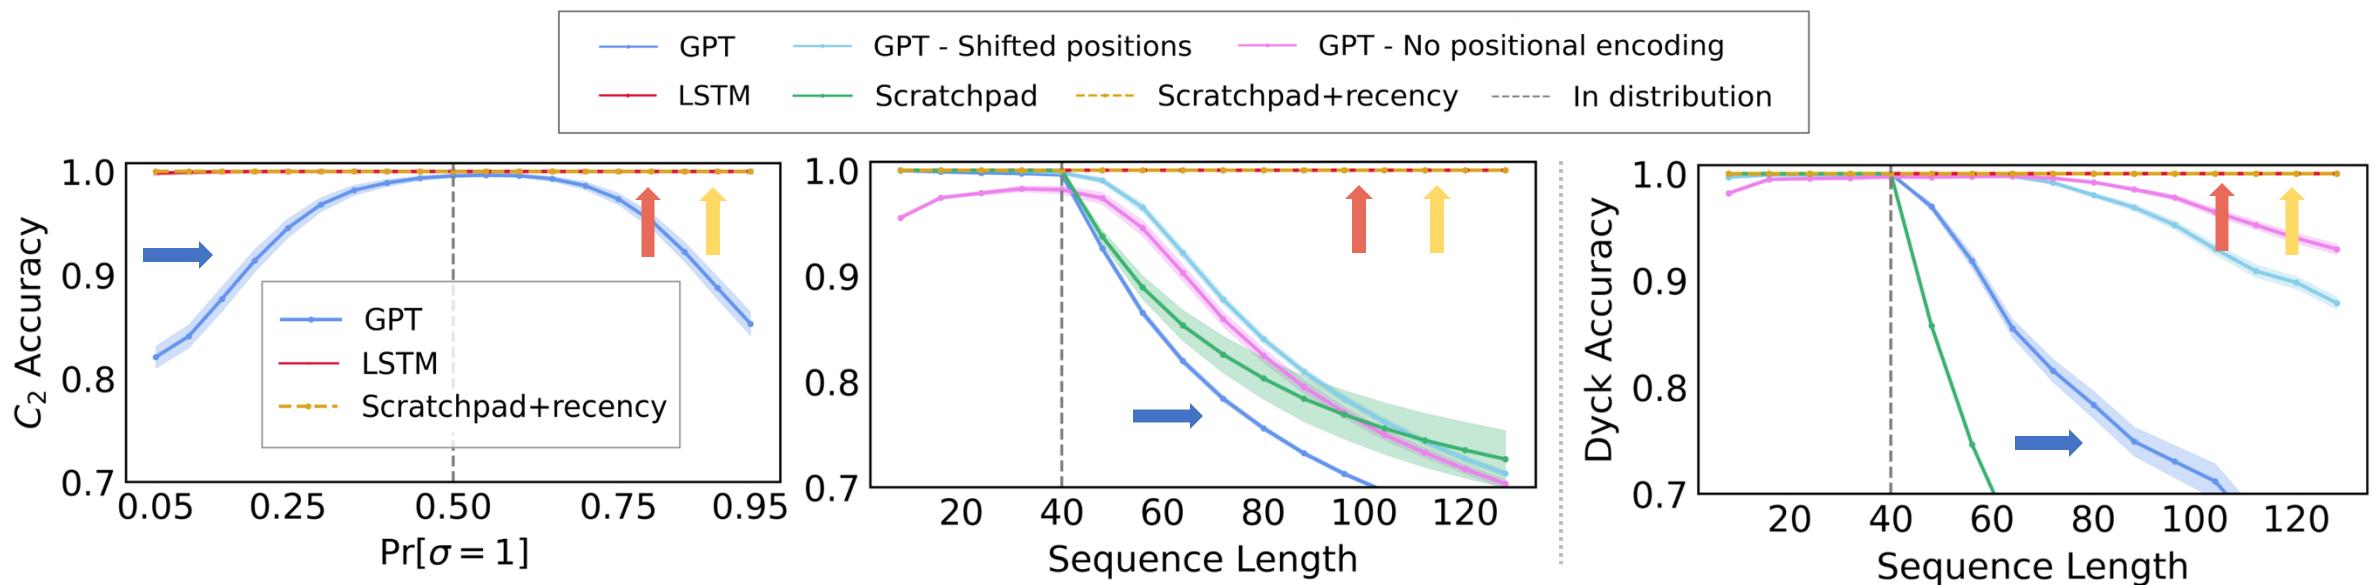
- *Setup:* train with previous states as inputs;



The recurrent mode of Transformers

Q3: Solutions? Guiding Transformers to learn recurrent solutions.

- *Setup:* train with previous states as inputs; eval at diff distributions or lengths.



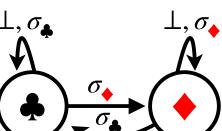
Takeaway: Transformers turned recurrent with **scratchpad** [Nye et al. 22] + **recency bias** [Press et al. 21] → Open: *Can we learn shortcuts that generalize?*

Takeaway

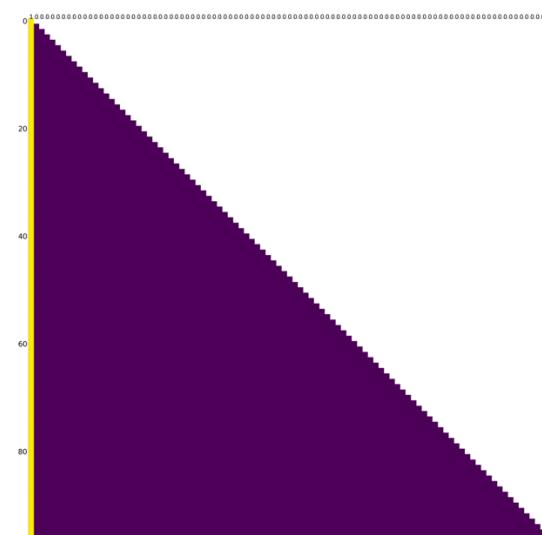
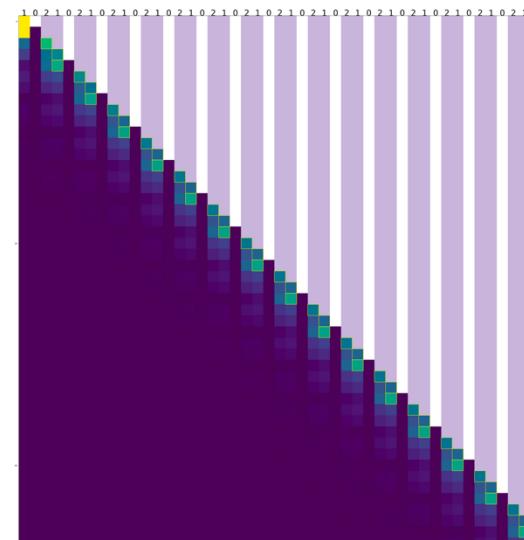
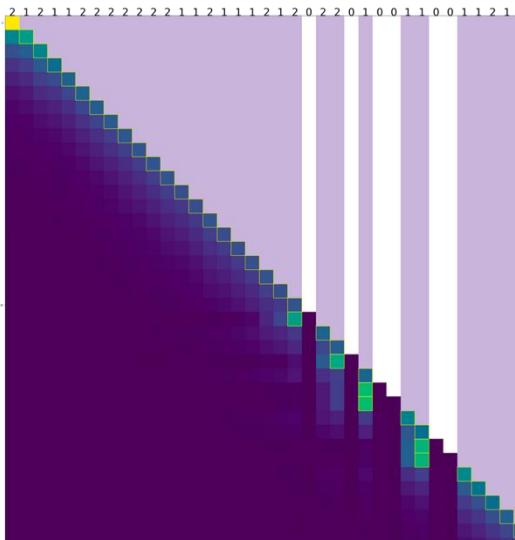
- **Positive:** Transformers with SGD can find good in-distribution solutions.
 - The trend roughly matches our beliefs on the difficulty of the tasks.
- **Challenges:** shortcuts are **less robust** than recurrent methods
 - Struggle at OOD generalization / changes in training setups.
 - Open question: How to thoroughly test for generalization?
 - Fix: made recurrent by scratchpad + recency bias.
- **Mechanistic understanding**
 - Gridworld: boundary detection.
 - Parity: evidence for implementing $(\sum_t \sigma_t) \bmod 2$.

Discussions: Interpretability and generalization

How to test for generalization? ∞ unseen inputs -- how much do we need to probe?

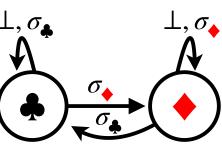
e.g. Flipflop  : train with $p(0) = 0.5, p(1) = p(2) = \frac{1-p(0)}{2}$.

- Test: $p(0) \in \{0.1, 0.2, \dots, 0.9\} \times (100k \text{ len-100 seqs})$: 100% accuracy

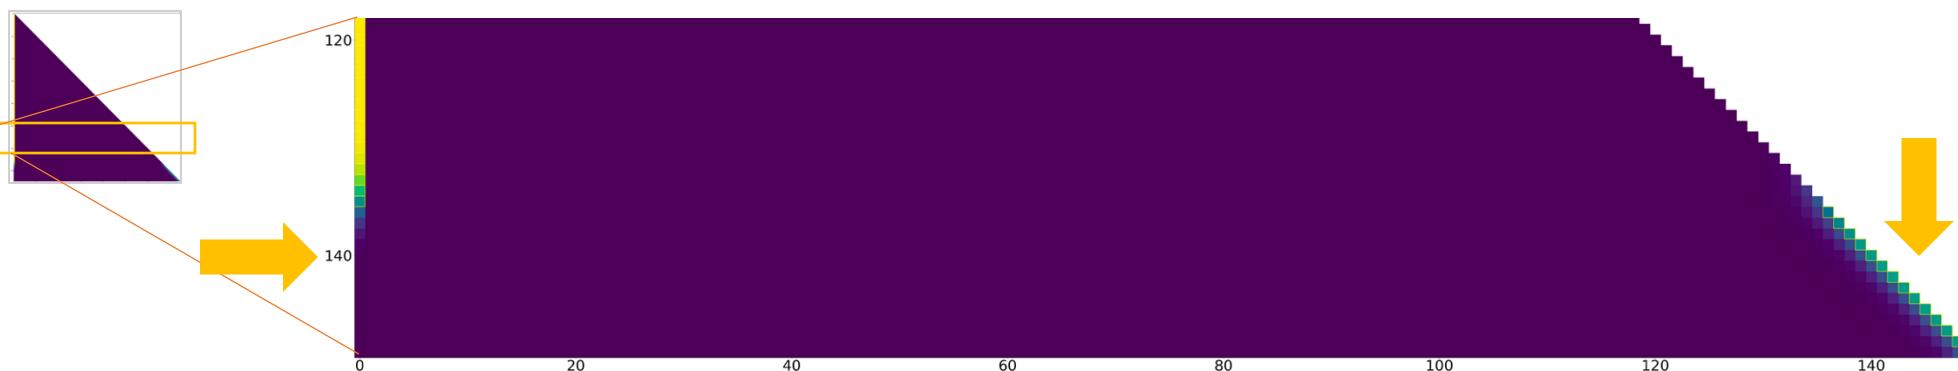


Discussions: Interpretability and generalization

How to test for generalization? ∞ unseen inputs -- how much do we need to probe?

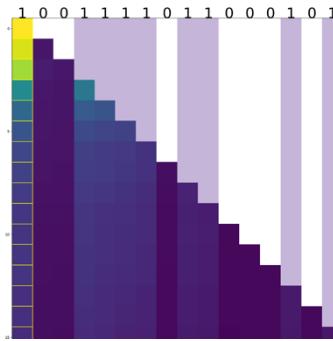
e.g. Flipflop  : train with $p(0) = 0.5, p(1) = p(2) = \frac{1-p(0)}{2}$.

- Test: 100% for $p(0) \in \{0.1, 0.2, \dots, 0.9\}$, for 100k len-100 seqs per $p(0)$.
- $[1,0,0,0, \dots, 0]$ is the hardest sample -- long-term dependency solved! 😊
- Not quite: failing on length generalization. 😞

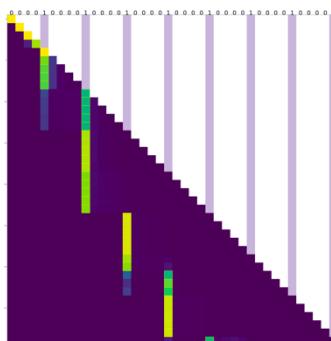


Discussions: (mechanistic) interpretability

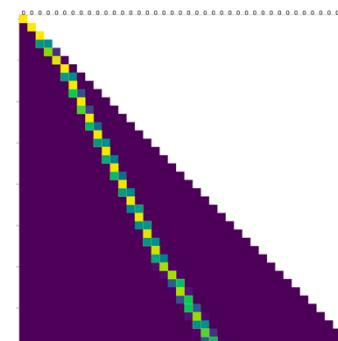
- Parity: what if not “sensible”? Variance (across randomness) in the patterns?



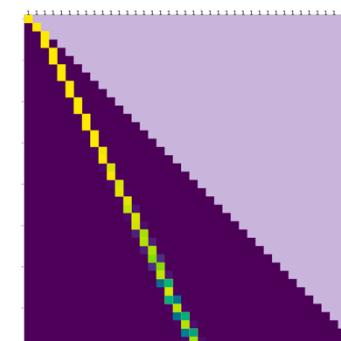
attend to 1s only



$[0,0,0,1]*10$

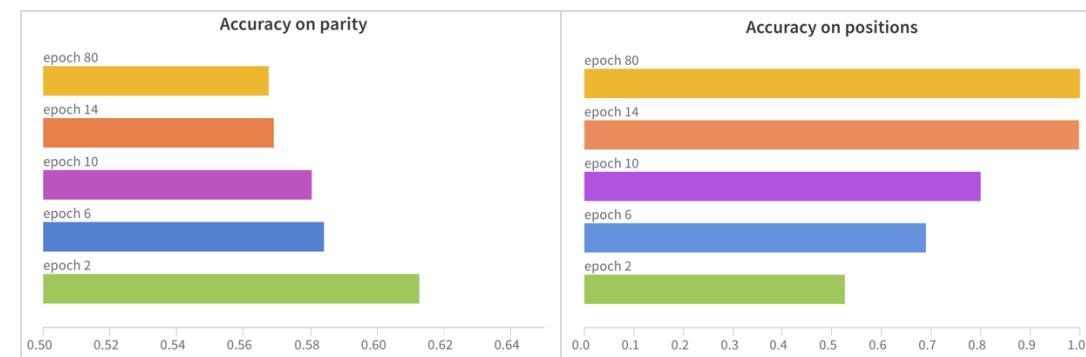
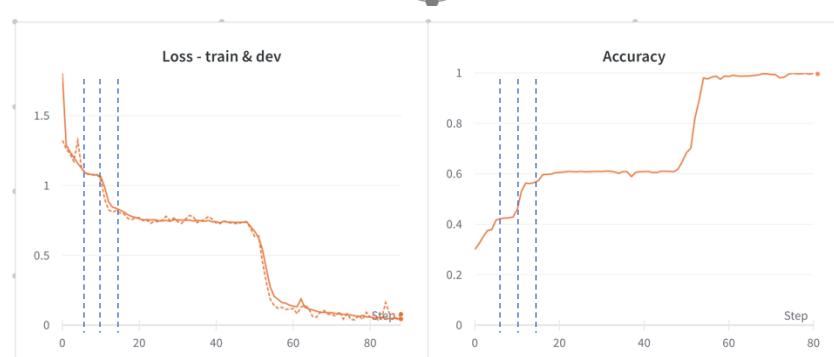


all-0s



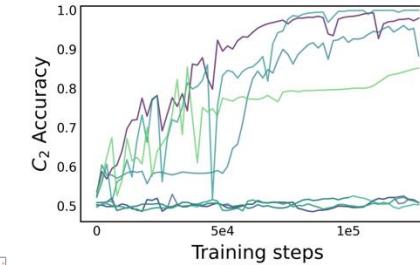
all-1s

- Dihedral group

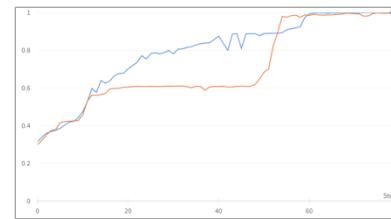


Discussions: optimization

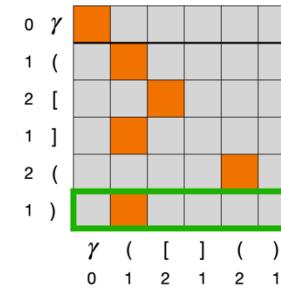
- Improve / stabilize training?
 - e.g. Parity: without positional encoding? With 1 layer?



- How to find (hidden) progress measures?
 - [Barak 22](#), [Nanda 23](#)



- What solutions are preferred, e.g. among shortcuts?
 - Dyck: sparse in theory ([Yao 20](#)), uniform in practice.
 - Are some solutions *better* than others?
- Modified architecture: parity: no issue with mod if using $\sin(\sum \sigma_t)$.
→ How to optimize with sinusoidal activation?



Discussions

- Generalization
 - Is it possible to have *guarantees* via tests (e.g. a “complete” test set)?
 - Is interpretability required for true generalization?
- Synthetic & real
 - *Why synthetic*: clean setup + cheaper + infinite supply of data
 - Decomposition → understand basic units first
 - Easier to formalize & understand & diagnose & fix
 - *Why not synthetic*: not directly useful
 - What’s missing: factual aspects? “Complex” structures?
 - *Goal*: Transfer insights to code / math / languages.
- Shortcut & recurrent: best of both worlds? ([RWKV](#)?)

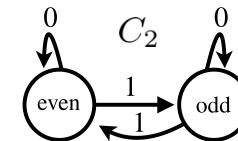
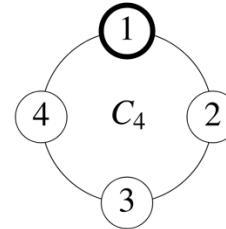
Transformers Learn Shortcuts to Automata

TL;DR: $o(T)$ -depth Transformers can simulate \mathcal{A} , in theory and practice

- **Theory:** for any length T , Transformers can learn:
 - $O(\log T)$ -layer for *all \mathcal{A}* – parallel prefix computation
 - Non-solvable \mathcal{A} : matching lower bound (Barrington)
 - $O(|Q|^2 \log |Q|)$ -layer for all *solvable \mathcal{A}* – factorization with Krohn-Rhodes
 - $O(1)$ -layer for *gridworld* – boundary detector
- **Empirical study:** shortcuts can be found but are *brittle* OOD.
 - Fix: make Transformers recurrent (scratchpad + recency bias).
- **Future work:** understanding and improving the model?

Appendix

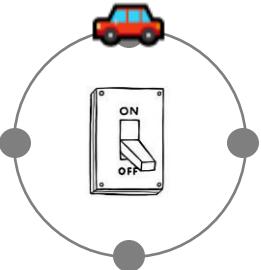
Theorem 2: the glue



Direct product \times , e.g. 🦓 $C_4 \times C_2$

Two **independent** groups

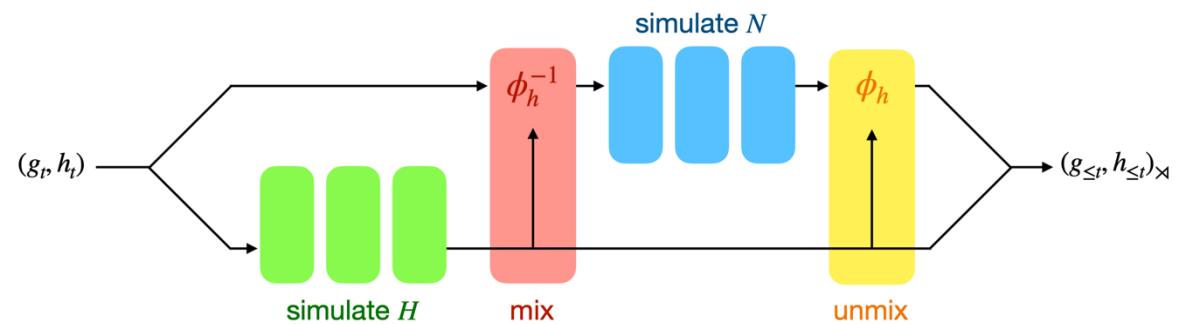
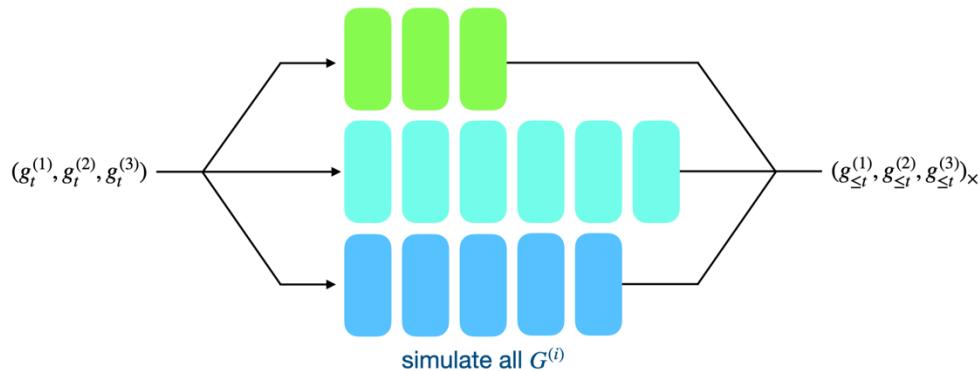
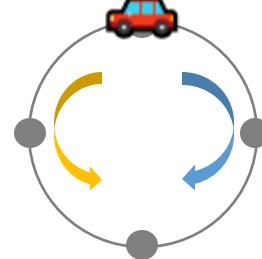
- $(g_1, h_1) \cdot (g_2, h_2) = (g_1g_2, h_1h_2)$
- e.g. car + a light switch



Semidirect product \rtimes , e.g. 🦀 $D_8 \cong C_4 \rtimes C_2$

Two **interacting** groups

- $(g_1, h_1) \cdot (g_2, h_2) = (g_1h_2g_2h_2^{-1}, h_1h_2)$
- e.g. car + direction toggle

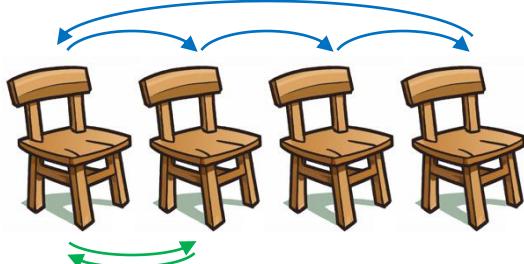


What about semigroups?

$\mathcal{T}(\mathcal{A}) := \{\delta(\cdot, \sigma) : \sigma \in \Sigma\}$ under composition

More complicated: **rank collapses**.

n -player musical chairs



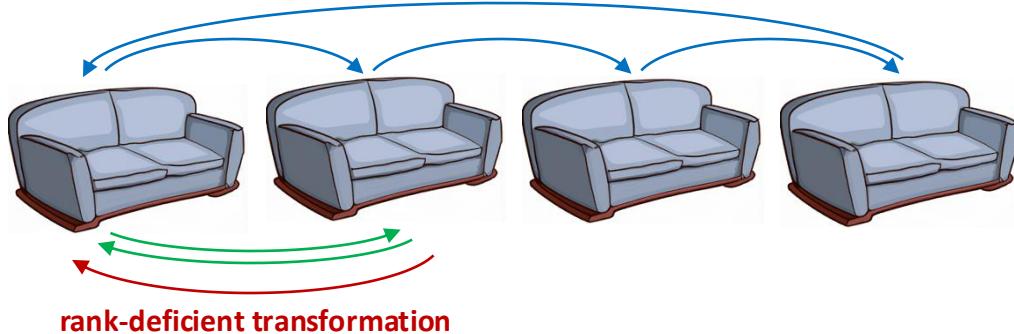
$$Q = \{\text{positions of } n \text{ players}\}$$

$$\Sigma = \{ \text{cycle, swap} \}$$

$$\begin{bmatrix} & & 1 \\ 1 & 1 & \end{bmatrix} \quad \begin{bmatrix} 1 & 1 \\ & 1 \end{bmatrix}$$

$\mathcal{T}(\mathcal{A}) = S_n$: all $n!$ permutations on $[n]$

n -player musical sofas



$$Q = \{\text{positions of } n \text{ players}\}$$

$$\Sigma = \{ \text{cycle, swap, merge} \}$$

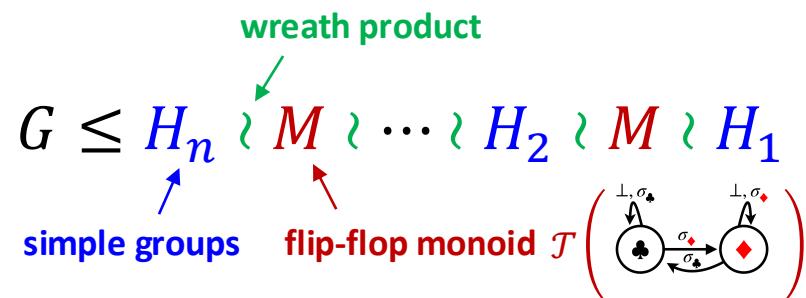
$$\begin{bmatrix} & & 1 \\ 1 & 1 & \end{bmatrix} \quad \begin{bmatrix} 1 & 1 \\ & 1 \end{bmatrix} \quad \begin{bmatrix} 1 & 1 \\ & 1 \end{bmatrix}$$

$\mathcal{T}(\mathcal{A}) = T_n$: all n^n functions $[n] \rightarrow [n]$

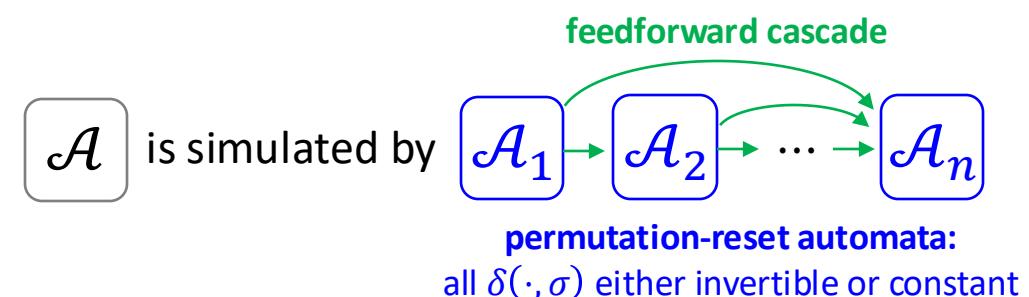
Prime factorizations of algebraic objects

- Universal structure theorems in abstract algebra:
 - Integers: $N = p_1 \cdot p_2 \cdots p_n$, prime numbers p_i [Euclid ~300 BC]
 - Groups: $G \triangleright H_n \triangleright \cdots \triangleright H_1$, simple groups H_{i+1}/H_i [Jordan & Hölder ~1880]
 - Semigroups: do we know anything if we're only given associativity?
- Yes! [Krohn & Rhodes '65]

Krohn-Rhodes theorem (semigroups)



Krohn-Rhodes theorem (semiautomata)



Factorization: from integers to groups

$$8 = 2 \times 2 \times 2$$

- Why groups get complicated: **combinatorial explosion**



C_8 : mod-8 addition



$E_8 \cong C_2 \times C_2 \times C_2$: 3-bit vectors under XOR



$C_4 \times C_2$: non-interacting mod-4 & parity



$D_8 \cong C_4 \rtimes C_2$: rotations/reflections of a square



Q_8 : multiplication of unit quaternions

} **non-abelian:** $gh \neq hg$

- **Finite group theory:** classical toolbox for understanding symmetries

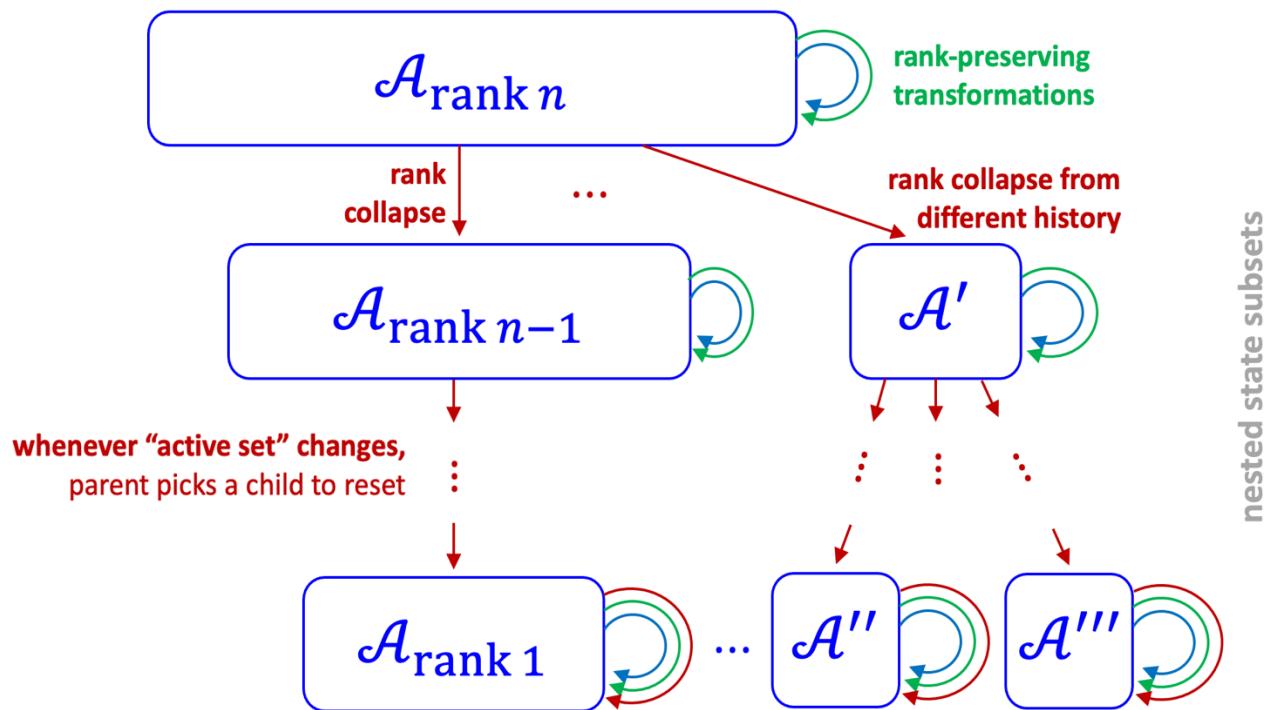
$$C_8, E_8, C_4 \times C_2, D_8, Q_8 \leq (C_2 \wr C_2) \wr C_2$$

Jordan-Hölder factors (simple groups)

Krasner-Kaloujnine embedding (wreath product)

Krohn-Rhodes intuitions

Tracking rank collapses (*holonomy decomposition*)

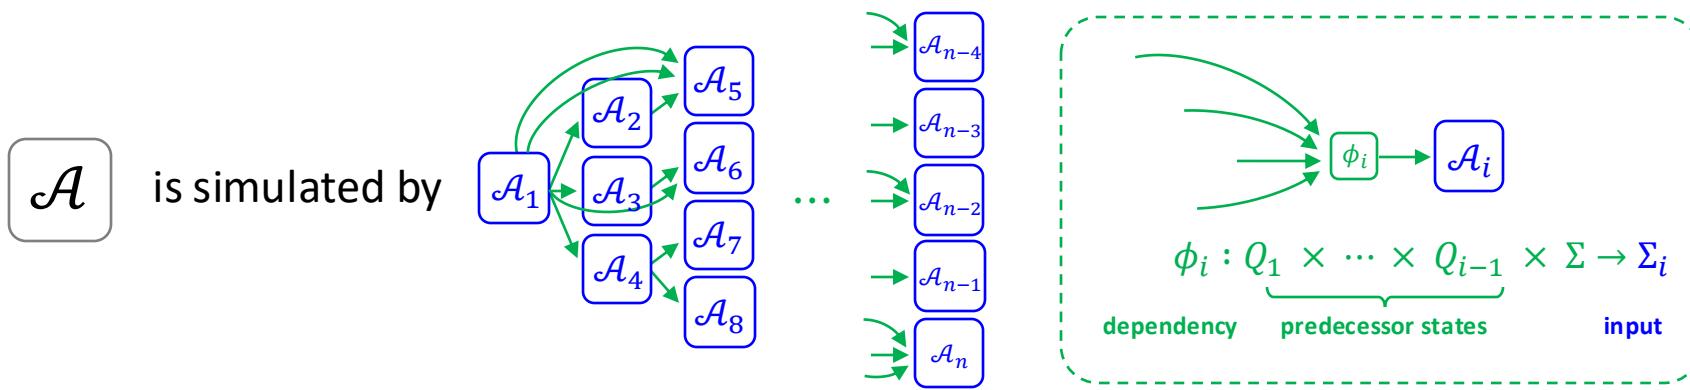


Number of layers:

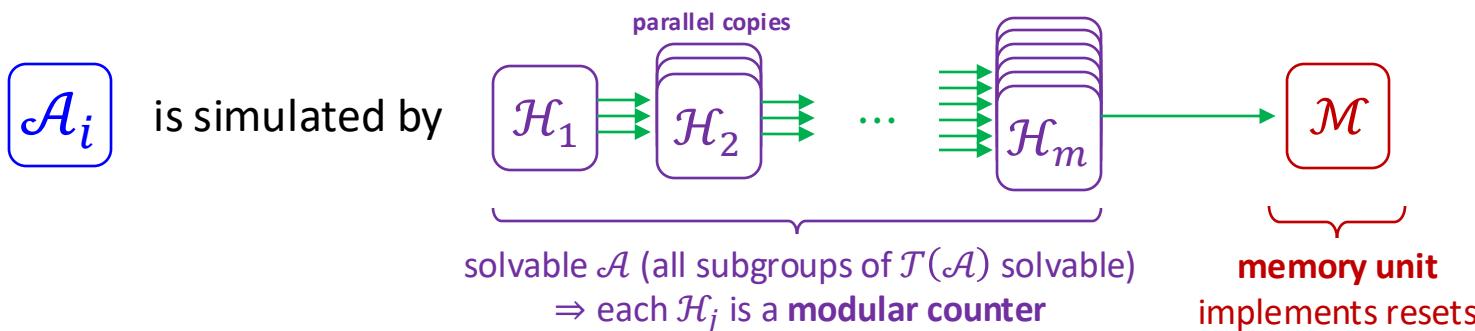
- Solvable groups: $O(\log|G|)$
 - mod counter
- Permutation-reset semiautomaton:
 $O(\log|G|) + 2 \leq O(|Q|\log|Q|)$.
 - mod counter + memory unit
- Semiautomaton: $\leq |Q|$ levels of the above.

Proof of Krohn-Rhodes: key intuitions

- Holonomy decomposition “compiles” to a cascade

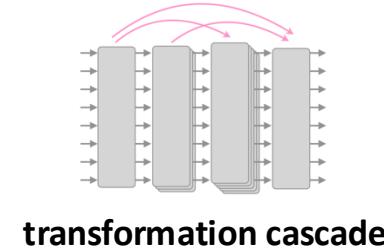
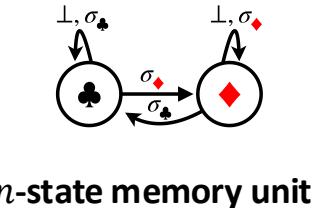
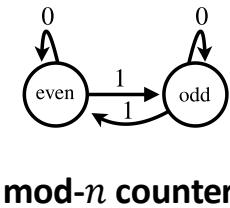


- Permutation-reset semiautomata factorize further



Implications for simulating automata

- Krohn-Rhodes: all \mathcal{A} decompose into permutation-reset $\{\mathcal{A}_i\}$
- Jordan-Hölder: all \mathcal{A}_i decompose into simple group machines $\{\mathcal{H}_i\}$
- Recall: when can we find a solution like $\Sigma \sigma_{1:t} \bmod 2$ for parity?
- Answer: whenever the “*atoms*” and “*glue*” are implementable

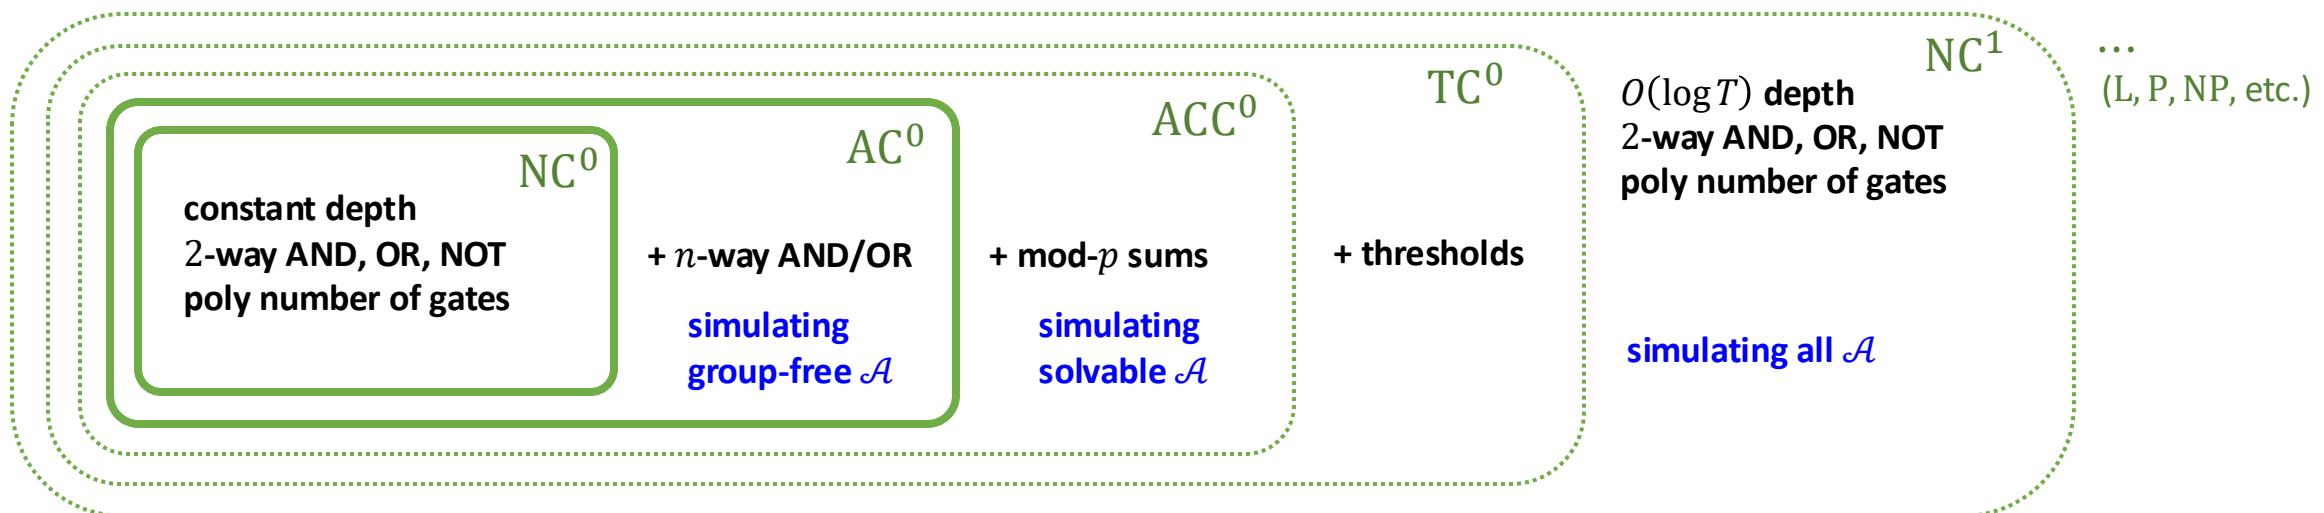


} sufficient when all groups within $\mathcal{T}(\mathcal{A})$ are **solvable**

- Need circuit complexity to formalize the question...

Quantifying efficient parallel circuits

- Goal: formalize “*Krohn-Rhodes implies efficient simulation*”
- Low-depth parallel algorithms are best captured by **circuit complexity**

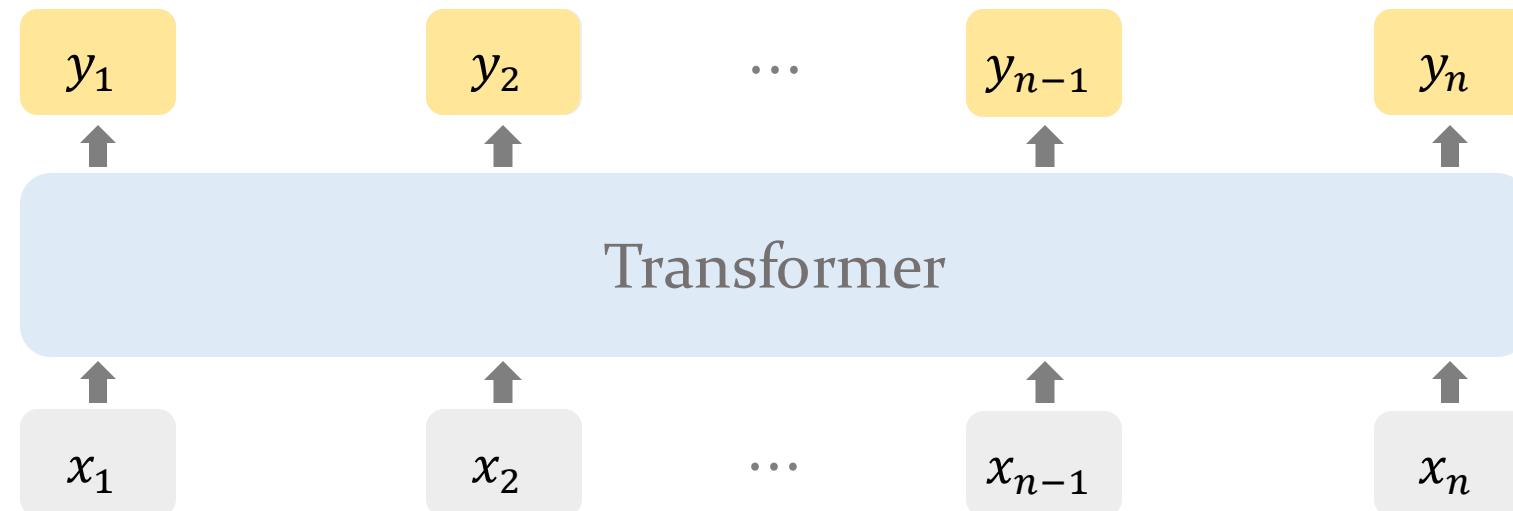


Embarrassingly open: are any of the \square proper? $ACC^0 = ?$ $NP = ?$



Scratchpads (modification of Nye et al. 21)

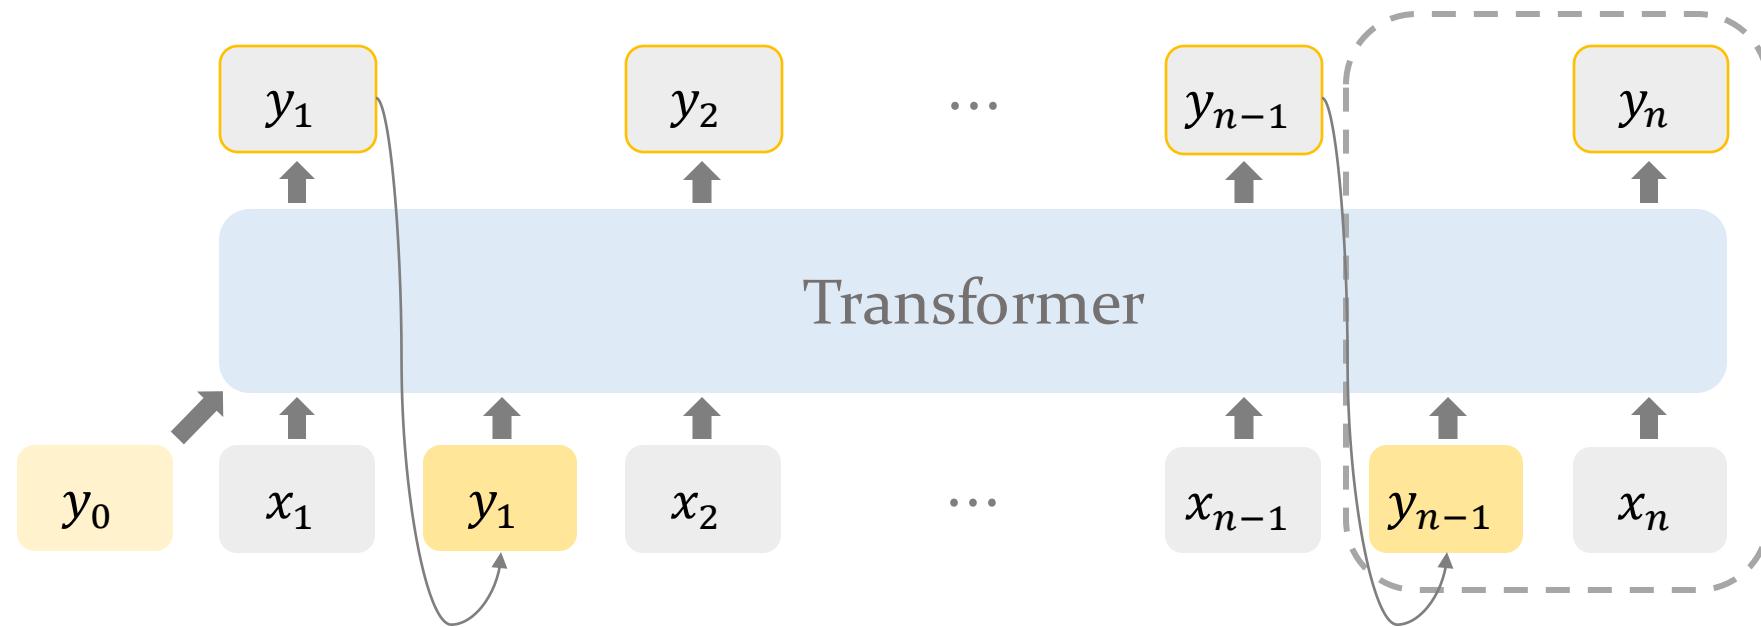
Idea: make the model **recurrent** with scratchpad (~buffer): explicitly modeling states.



$$\text{want: } y_i = x_i \oplus y_{i-1}$$

Scratchpads (modification of Nye et al. 21)

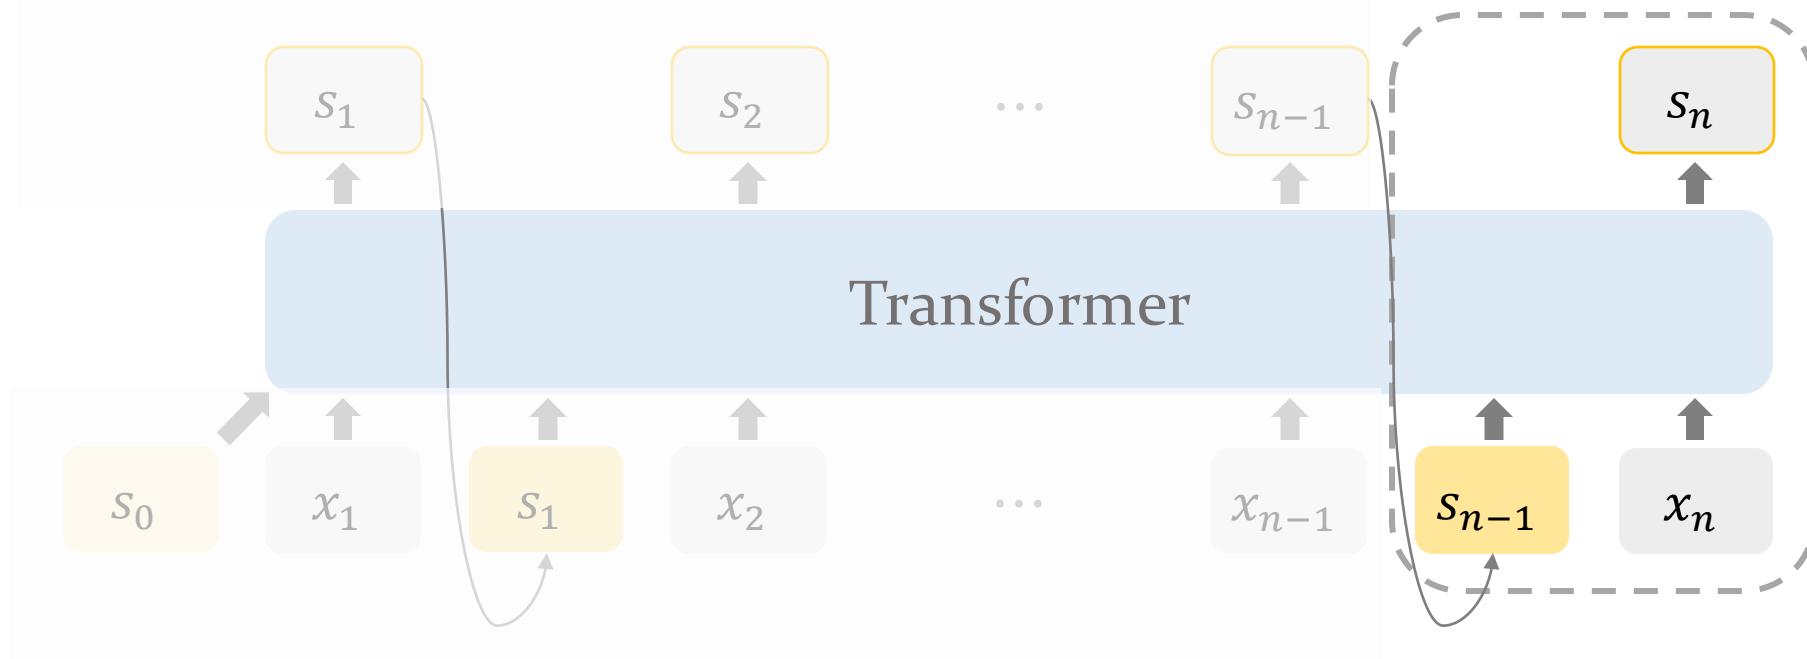
Idea: make the model recurrent with **scratchpad** (~buffer): explicitly modeling states



$$\text{want: } y_i = x_i \oplus y_{i-1}$$

Scratchpads (modification of Nye et al. 21)

Idea: make the model recurrent with **scratchpad** (~buffer): explicitly modeling states



$$\text{want: } s_i := y_i = x_i \oplus s_{i-1}$$